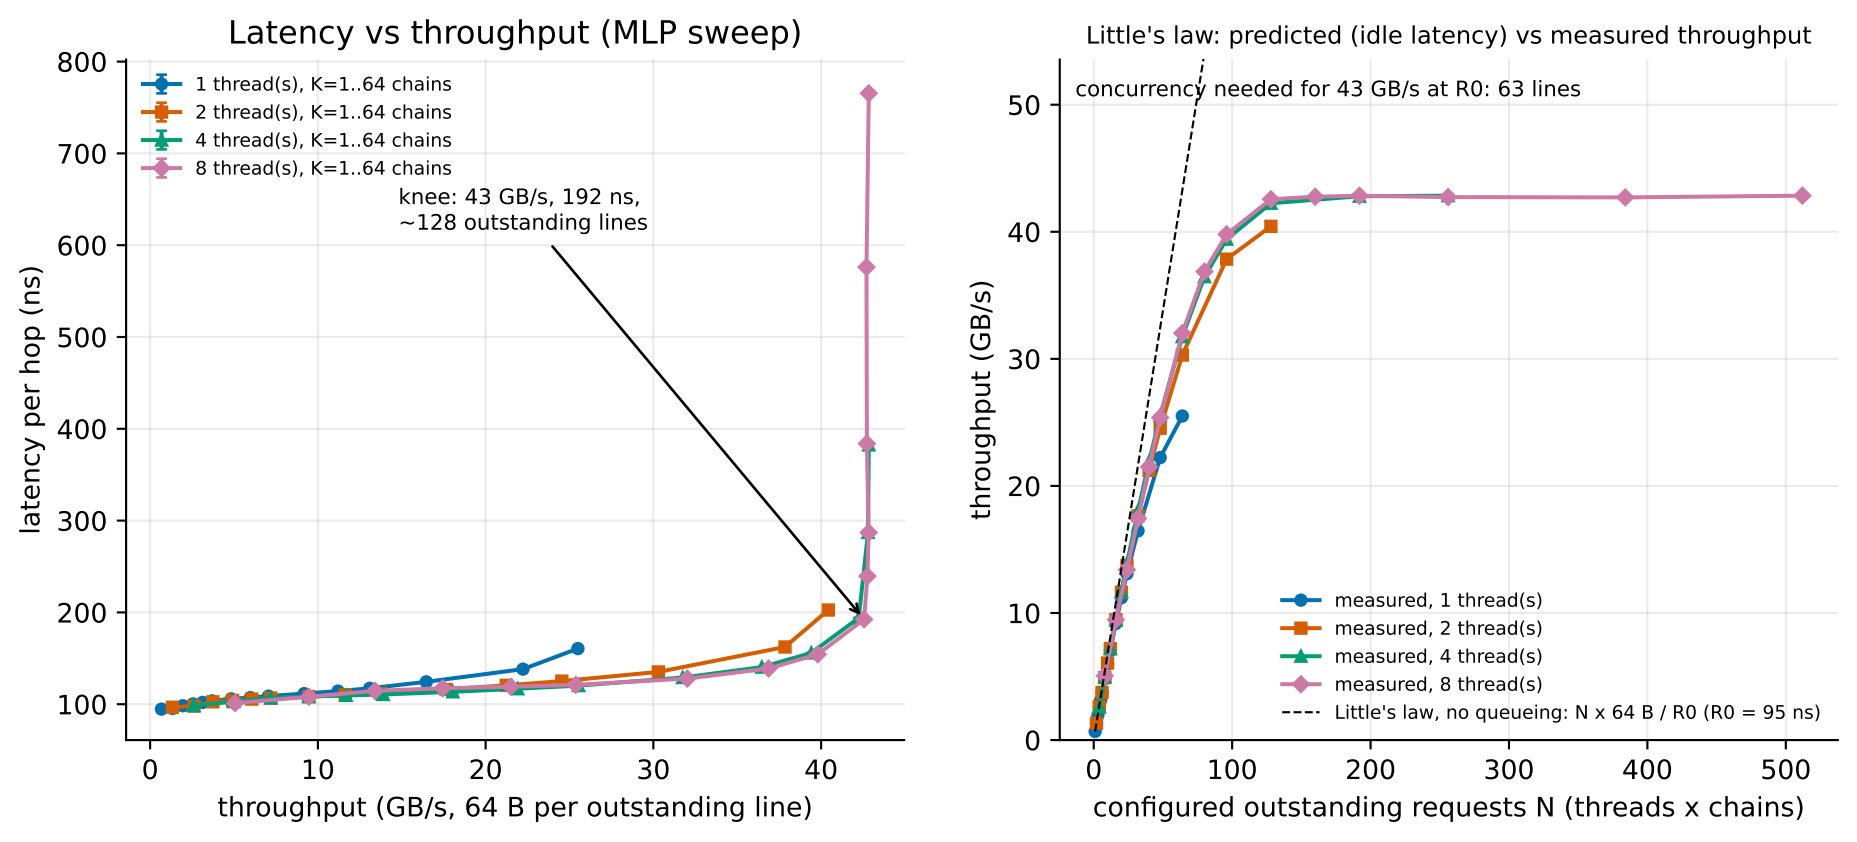
Latency vs throughput (MLP sweep)
1 thread(s), K=1..64 chains
2 thread(s), K=1..64 chains
4 thread(s), K=1..64 chains
8 thread(s), K=1..64 chains
knee: 43 GB/s, 192 ns,
~128 outstanding lines
latency per hop (ns)
throughput (GB/s, 64 B per outstanding line)
Little's law: predicted (idle latency) vs measured throughput
concurrency needed for 43 GB/s at R0: 63 lines
measured, 1 thread(s)
measured, 2 thread(s)
measured, 4 thread(s)
measured, 8 thread(s)
Little's law, no queueing: N x 64 B / R0 (R0 = 95 ns)
throughput (GB/s)
configured outstanding requests N (threads x chains)